

# Summarising and Plotting Data in R

## Summarising Data

Glenn Williams

University of Sunderland

2021-10-27 (updated: 2021-12-05)



# Summarising Data

Individual data points are some **abstraction away from reality**, essentially removing some detail.

- Even with trial-level data, we will **summarise responses** in terms of reaction times or correct responses.
- This removes a lot of the details of the real experience, but allows us to **see patterns**.
- Similarly, we can summarise data from many participants in a number of ways to help communicate a pattern or idea.

We can present summaries of data as a **plot**, as a **table of descriptive statistics**, and make inferences and decisions based on summaries of the data by **modelling our data**.

# Descriptive Statistics

We often summarise data in terms of measures of **central tendency** and **dispersion**.

- **Central tendency:** What is the average score? This can be determined in a few ways.
  - **Mean:** used to describe data that are normally distributed. Can be misleading when outliers are present.
  - **Median:** used to describe data that are skewed in some way away from a normal distribution. Suppresses the impact of outliers.

```
noskew <- c(10, 20, 15, 20, 22)
mean(noskew)
```

```
## [1] 17.4
```

```
skew <- c(10, 20, 15, 20, 2000)
mean(skew)
```

```
## [1] 413
```

```
median(noskew)
```

```
## [1] 20
```

```
median(skew)
```

```
## [1] 20
```

# Descriptive Statistics

- **Dispersion:** The spread of scores (e.g. for individuals) around an average.
  - **Standard Deviation:** the spread of the data from the sample mean. With a normal distribution approximately 1, 2, and 3 SDs capture approximately 68, 95, and 99.7% of the data.
  - **Interquartile Range:** Rank orders data into four equal parts; the first region – or quartile – (i.e. 25%), the median (i.e. 50%), and the middle part of the third quartile (i.e. 75%). Used for non-normal data.

```
# low variability in scores = low SD
low_dispersion <- c(10, 12, 8, 9, 11, 10, 10.5, 9, 11.5)
sd(low_dispersion)
```

```
## [1] 1.293681
```

```
# lots of variability in scores = high SD
high_dispersion <- c(10, 1000, 89, -400, 90, 880, 0)
sd(high_dispersion)
```

```
## [1] 508.6885
```

# Why Make these in R?

- **Scales up easily to new data:** If you write this code to analyse data from 10 participants, you can instantly rerun it for millions.
- **Easier to spot mistakes** when you're writing the recipe.
- **Easier to fix mistakes** with a minor modification instead of repeating every step of the analysis by hand.
- **Repeatability:** If you do a task once in R, you can copy and paste code for a new study. This saves a lot of time in the long run.
- **Reproducibility:** Allows others (and future you) to check work and inspect methods at every step. You'll make science more reliable!

# Summarise

Let's take a look at the inbuilt `starwars` data set in the `tidyverse` package. There's a lot of data here! How might we summarise this?

```
data(starwars) # load the data from within the package
glimpse(starwars) # view it
```

```
## Rows: 87
## Columns: 14
## $ name      <chr> "Luke Skywalker", "C-3PO", "R2-D2", "Darth Vader", "Leia Or...
## $ height    <int> 172, 167, 96, 202, 150, 178, 165, 97, 183, 182, 188, 180, 2...
## $ mass      <dbl> 77.0, 75.0, 32.0, 136.0, 49.0, 120.0, 75.0, 32.0, 84.0, 77....
## $ hair_color <chr> "blond", NA, NA, "none", "brown", "brown, grey", "brown", N...
## $ skin_color <chr> "fair", "gold", "white, blue", "white", "light", "light", "...
## $ eye_color  <chr> "blue", "yellow", "red", "yellow", "brown", "blue", "blue",...
## $ birth_year <dbl> 19.0, 112.0, 33.0, 41.9, 19.0, 52.0, 47.0, NA, 24.0, 57.0, ...
## $ sex        <chr> "male", "none", "none", "male", "female", "male", "female",...
## $ gender     <chr> "masculine", "masculine", "masculine", "masculine", "femini...
## $ homeworld  <chr> "Tatooine", "Tatooine", "Naboo", "Tatooine", "Alderaan", "T...
## $ species    <chr> "Human", "Droid", "Droid", "Human", "Human", "Human", "Huma...
## $ films      <list> <"The Empire Strikes Back", "Revenge of the Sith", "Return...
## $ vehicles   <list> <"Snowspeeder", "Imperial Speeder Bike">, <>, <>, <>, "Imp...
## $ starships  <list> <"X-wing", "Imperial shuttle">, <>, <>, "TIE Advanced x1",...
```

# Summarise

`summarise()` collapses across all observations in your data set to produce a single row of data.

- `summarise()` takes two main arguments:
  - 1: what is the **data** set you want to summarise?
  - 2: additional information specifying **how you want to summarise it**.

Let's summarise the mean height of the Star Wars characters:

```
summarise(starwars, mean_height = mean(height))
```

```
## # A tibble: 1 x 1
##   mean_height
##         <dbl>
## 1          NA
```

Oops, we got NA! Why? R by default produces NA for any summary of data containing NA values. (You can't average something that isn't there.)



# Summarise

What should we do? We can make a new data set without NAs, or tell `mean()` to ignore them.

```
summarise(starwars, mean_height = mean(height, na.rm = TRUE))
```

```
## # A tibble: 1 x 1
##   mean_height
##         <dbl>
## 1         174.
```

We set the argument **na.rm** in `mean()` to **TRUE**, meaning "Should R remove NAs? TRUE (yes)".

# Many Summaries

- What if we want to get **different types of summaries**? Imagine we want the count, mean, and standard deviation of height?
- We also can tell R to generate these summaries on the data by first **dropping any observations with missing values (NA)** in height.

```
summarise(  
  drop_na(starwars, height),  
  n = n(),  
  sd = sd(height),  
  mean_height = mean(height)  
)
```

```
## # A tibble: 1 x 3  
##       n      sd mean_height  
##   <int> <dbl>      <dbl>  
## 1     81  34.8      174.
```

- How might we interpret this? We have 81 measured heights of characters. The mean height is 174.36cm, while the standard deviation is 34.77.
- **What's the SD mean?** With normally distributed data 95% of values fall within  $\pm 1.96$  SD of the mean!

# Grouping

- As we can see, the data set contains **very many species** of creates from the Star Wars universe. What if we want summaries of all of them?
- We can add an argument to our data in summarise() called group\_by().
- Within group\_by(), we specify our data but also give a column by which to **group the summaries**.

```
summarise(  
  group_by(starwars, species),  
  mean_height = mean(height, na.rm = TRUE)  
)
```

```
## # A tibble: 38 x 2  
##   species  mean_height  
##   <chr>      <dbl>  
## 1 Aleena          79  
## 2 Besalisk       198  
## 3 Cerean         198  
## # ... with 35 more rows
```

# Getting Complicated

We can combine tidyverse functions together. We could filter before summarising the data, and even make multiple grouped summaries.

```
summarise(  
  group_by(filter(starwars, mass < 400), species),  
  mean_height = mean(height, na.rm = TRUE),  
  sd_height = sd(height, na.rm = TRUE),  
  n = n()  
)
```

```
## # A tibble: 31 x 4  
##   species mean_height sd_height    n  
##   <chr>      <dbl>      <dbl> <int>  
## 1 Aleena          79         NA     1  
## 2 Besalisk        198         NA     1  
## 3 Cerean          198         NA     1  
## # ... with 28 more rows
```

# Interim Recap

We've learned...

- Why we summarise data and recapped on basic **descriptive statistics**.
- **Why we might use R** to summarise our data.
- How to pass **optional arguments** to functions (e.g. `na.rm` in `mean()` and `sd()`).
- How to make **separate summaries per group** using `group_by()`.
- How to combine tidyverse functions together to get the summaries we need.

# Analysing Data

# Some Background

Data analysis is surprisingly one of the easiest parts of working with R.

- Once your data is in the correct (long) format, analysis using any test is highly consistent.
- We rely on a formula interface like this:

```
DV ~ IV
```

- Our dependent variable/predicted variable goes to the left of the ~ (tilde), while our independent variables or predictors go to the right.
- After this we specify our data:

```
DV ~ IV, data
```

We then apply a function to our formula which is the name of our test. There's some minor options we can choose within tests, but that's pretty much it!

# Correlations

## The Data

Let's check out the **starwars** data set again. We'll use this for our tests.

```
starwars <- starwars %>% filter(mass < 500)
```

We will use the height and mass columns, looking at whether mass is associated with height.

```
## Rows: 58
## Columns: 14
## $ name      <chr> "Luke Skywalker", "C-3P0", "R2-D2", "Darth Vader", "Leia Or...
## $ height    <int> 172, 167, 96, 202, 150, 178, 165, 97, 183, 182, 188, 228, 1...
## $ mass      <dbl> 77.0, 75.0, 32.0, 136.0, 49.0, 120.0, 75.0, 32.0, 84.0, 77....
## $ hair_color <chr> "blond", NA, NA, "none", "brown", "brown, grey", "brown", N...
## $ skin_color <chr> "fair", "gold", "white, blue", "white", "light", "light", "...
## $ eye_color  <chr> "blue", "yellow", "red", "yellow", "brown", "blue", "blue",...
## $ birth_year <dbl> 19.0, 112.0, 33.0, 41.9, 19.0, 52.0, 47.0, NA, 24.0, 57.0, ...
## $ sex        <chr> "male", "none", "none", "male", "female", "male", "female",...
## $ gender     <chr> "masculine", "masculine", "masculine", "masculine", "femini...
## $ homeworld  <chr> "Tatooine", "Tatooine", "Naboo", "Tatooine", "Alderaan", "T...
## $ species    <chr> "Human", "Droid", "Droid", "Human", "Human", "Human", "Huma..."
```



# Correlation

- Here, we aren't predicting any one variable from the other, so both variables go to the right of the tilde.
- We add multiple variables with a +.
- We choose the type of correlation we want (e.g. Pearson, Spearman) with the method.

```
cor.test(~ height + mass, starwars, method = "pearson")
```

```
##  
##      Pearson's product-moment correlation  
##  
## data:  height and mass  
## t = 8.7853, df = 56, p-value = 4.018e-12  
## alternative hypothesis: true correlation is not equal to 0  
## 95 percent confidence interval:  
##  0.6260700 0.8520232  
## sample estimates:  
##           cor  
## 0.7612612
```

# Tests of Difference

- We'll use some different data here on out.
- Let's assume this data looks at giving people a placebo or drug, and tests the effect of that drug at two different time points.
- We care about improvements in reaction times.

```
mixed_data <- read_csv(here("data", "mixed_factorial.csv"))  
head(mixed_data)
```

```
## # A tibble: 6 x 4  
##   id    drug    time      rt  
##   <chr> <chr>   <chr>   <dbl>  
## 1 S001 control daylater  431.  
## 2 S001 control monthlater 421.  
## 3 S002 control daylater  372.  
## 4 S002 control monthlater 350.  
## 5 S003 control daylater  393.  
## 6 S003 control monthlater 368.
```

# t-tests

## One-sample t-test

- We have only one variable here, so we don't even need a formula.
- We compare the mean of this variable against a specified baseline mean (here 400).

```
t.test(mixed_data$rt, mu = 400)
```

```
##  
##      One Sample t-test  
##  
## data:  mixed_data$rt  
## t = -9.5311, df = 479, p-value < 2.2e-16  
## alternative hypothesis: true mean is not equal to 400  
## 95 percent confidence interval:  
##  376.3267 384.4193  
## sample estimates:  
## mean of x  
##    380.373
```

# t-tests

## Independent-samples t-test

- Do reaction times vary depending on the drug given to participants?
- We test reaction times predicted by drug, with a regular t-test where variances are assumed to be equal (`var.equal = TRUE`).

```
t.test(rt ~ drug, mixed_data, var.equal = TRUE)
```

```
##  
##      Two Sample t-test  
##  
## data:  rt by drug  
## t = 3.732, df = 478, p-value = 0.0002128  
## alternative hypothesis: true difference in means is not equal to 0  
## 95 percent confidence interval:  
##    7.18135 23.15251  
## sample estimates:  
##    mean in group control mean in group treatment  
##           387.9564           372.7895
```

# t-tests

## Paired t-test

- Do reaction times vary over time (i.e. practice)?
- We test reaction times predicted by the time of testing. This is a paired test (`paired = TRUE`) and a regular t-test where variances are assumed to be equal (`var.equal = TRUE`).

```
t.test(rt ~ time, mixed_data, paired = TRUE, var.equal = TRUE)
```

```
##  
##      Paired t-test  
##  
## data:  rt by time  
## t = 18.348, df = 239, p-value < 2.2e-16  
## alternative hypothesis: true difference in means is not equal to 0  
## 95 percent confidence interval:  
##  49.10644 60.91970  
## sample estimates:  
## mean of the differences  
##                55.01307
```

# One-way ANOVA

## Between-subjects

- What if we had **more than two groups** for the drug condition? We use an ANOVA.
- We simply change the test function to `aov()` (**A**nalysis **O**f **V**ariance)
- We need to summarise the model results here to get a regular ANOVA output.

```
summary(aov(rt ~ drug, mixed_data))
```

```
##              Df Sum Sq Mean Sq F value    Pr(>F)
## drug           1  27604    27604   13.93 0.000213 ***
## Residuals    478 947379     1982
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

# One-way ANOVA

## Within-subjects

- What if we have more than two groups and a **within-subjects design**?
- We do the same as before, but need to add an **Error term** to the formula. This states that we adjust our errors to account for the fact scores in each group belong to the same participant (i.e. **id** in our data).

```
summary(aov(rt ~ time + Error(id), mixed_data))
```

```
##
## Error: id
##           Df Sum Sq Mean Sq F value Pr(>F)
## Residuals 239 353969    1481
##
## Error: Within
##           Df Sum Sq Mean Sq F value Pr(>F)
## time       1 363173  363173   336.6 <2e-16 ***
## Residuals 239 257842    1079
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

# Two-Way ANOVA

## Mixed

```
summary(aov(rt ~ time * drug + Error(id), mixed_data))
```

```
##
## Error: id
##           Df Sum Sq Mean Sq F value    Pr(>F)
## drug         1  27604    27604   20.13 1.13e-05 ***
## Residuals  238 326365     1371
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Error: Within
##           Df Sum Sq Mean Sq F value    Pr(>F)
## time         1 363173   363173   806.3 <2e-16 ***
## time:drug     1 150636   150636   334.4 <2e-16 ***
## Residuals  238 107206     450
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```



# Throw Away the Alphabet Soup

All of the statistical tests you know (e.g. *t*-tests, ANOVA, chi-square) are just extensions of the **general linear model**. This is the most important thing you can learn to use in statistics.

Learn the mean and *variance* of some measurement by using an additive combination of other measurements.

- The **geocentric model of applied statistics**: used wisely, can be useful. But we shouldn't read too much into the numbers produced. They're almost certainly wrong because we can't (and shouldn't) model all sources of variance.
- Predict a **linear relationship** between one or more variable(s) and a continuous (e.g. scale) dependent variable.
- Predictor variables can be continuous or categorical.

# Linear Regression

Takes the general form:

$$Y = \alpha + \beta X + e$$

- **Outcome**  $Y$  = intercept + (slope  $\times$   $X$ ) + residual error
- **Residuals**  $e$  = distance of observed values from predicted values
- *Note:* We do not fit a perfect model, hence the error term. This is a good thing, otherwise we are probably **overfitting** to our data; relying too much on our observed sample to draw inferences.

# Linear Regression

Takes the general form:

$$Y = \alpha + \beta X + e$$

- The **intercept**,  $\alpha$ , is usually the point on the y-axis at the lowest value of X (usually 0).
- The **slope**,  $\beta$ , corresponds to how much Y increases by for every increment in X.
- The **error**,  $e$ , corresponds to a constant by which to add to our estimates accounting for additional variation from other sources that we do not model.

# Linear Regression

Fit the model predicting height from weight from the starwars data.

$$Y = \alpha + \beta X + e$$

$$\text{height} = \text{intercept} + \text{slope} \times \text{mass} + \text{error}$$

```
starlm <- lm(height ~ mass, starwars)
summary(starlm)
```

```
##
## Call:
## lm(formula = height ~ mass, data = starwars)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -53.369  -6.816   2.042  13.851  44.719
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)  103.5133     8.5937  12.045  < 2e-16 ***
## mass          0.9327     0.1062   8.785 4.02e-12 ***
## ---
```

# Comparing tests we know...

## Correlation

```
broom::tidy(cor.test(~ height + mass, starwars, method = "pearson"))
```

```
## # A tibble: 1 x 8
##   estimate statistic  p.value parameter conf.low conf.high method      alternative
##   <dbl>      <dbl>    <dbl>      <int>    <dbl>    <dbl> <chr>      <chr>
## 1    0.761      8.79 4.02e-12         56    0.626    0.852 Pearson'... two.sided
```

```
broom::tidy(lm(height ~ mass, starwars, method = "pearson"))
```

```
## Warning in lm(height ~ mass, starwars, method = "pearson"): method = 'pearson'
## is not supported. Using 'qr'
```

```
## # A tibble: 2 x 5
##   term          estimate std.error statistic  p.value
##   <chr>          <dbl>    <dbl>    <dbl>    <dbl>
## 1 (Intercept)  104.        8.59     12.0 3.53e-17
## 2 mass          0.933     0.106     8.79 4.02e-12
```

# Comparing tests we know...

## t-tests

```
broom::tidy(t.test(rt ~ drug, mixed_data, var.equal = TRUE))
```

```
## # A tibble: 1 x 10
##   estimate estimate1 estimate2 statistic  p.value parameter conf.low conf.high
##   <dbl>     <dbl>     <dbl>     <dbl>    <dbl>     <dbl>     <dbl>     <dbl>
## 1    15.2     388.     373.      3.73 0.000213      478      7.18     23.2
## # ... with 2 more variables: method <chr>, alternative <chr>
```

```
broom::tidy(summary(lm(rt ~ drug, mixed_data)))
```

```
## # A tibble: 2 x 5
##   term          estimate std.error statistic  p.value
##   <chr>         <dbl>     <dbl>     <dbl>    <dbl>
## 1 (Intercept)    388.        2.87     135.      0
## 2 drugtreatment -15.2        4.06     -3.73 0.000213
```

- *t* statistics match exactly.

# Comparing tests we know...

## ANOVA

```
summary(aov(rt ~ drug, mixed_data))
```

```
##              Df Sum Sq Mean Sq F value    Pr(>F)
## drug           1  27604    27604    13.93 0.000213 ***
## Residuals     478 947379     1982
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

```
broom::tidy(lm(rt ~ drug, mixed_data))
```

```
## # A tibble: 2 x 5
##   term          estimate std.error statistic  p.value
##   <chr>          <dbl>     <dbl>     <dbl>    <dbl>
## 1 (Intercept)    388.         2.87     135.      0
## 2 drugtreatment -15.2         4.06     -3.73 0.000213
```

- $t$  to  $F$  is just  $t$  squared. So, 3.732 squared = 13.93...

# Bye!



*Effect sizes are easily handled by the `{effectsize}` package. Super-easy ANOVAs are done using the `{afex}` package.*